

SPARK BASED TOPICS KEYWORDS:

Spark Intro:

1. **Spark** : In-memory processing engine
2. **Why spark is fast**: Due to less I/O disc reads and writes
3. **RDD**: It is a data structure to store data in spark
4. **When RDD fails**: Using lineage graph we track which RDD failed and reprocess it
5. **Why RDD immutable** : As it has to be recovered after its failure and to track which RDD failed
6. Operations in spark: Transformation and Action
7. Transformation: Change data from one form to another, are lazy.
8. Action: Operations which processes the transformations, not lazy. creates DAG to remember sequence of steps.
9. Broadcast Variables: Data which is distributed to all the systems. Similar to map side join in hive
10. Accumulators: Shared copy in driver, executors can update but not read. Similar to counters in MR
11. MR before Yarn: Job tracker (scheduling & monitoring), task manager (manages tasks in its node)
12. Limitations of MR: Unable to add new clusters(scalable), resource under-utilization, only MR jobs handled
13. YARN: Resource manager(scheduling), application master(monitors & resource negotiation), node manager (manages tasks in its node)
14. Uberization: Tasks run on AM itself if they are very small

15. **Spark components**: Driver (gives location of executors) and executors(process data in memory)
16. **Client Mode**: Driver is at client side
17. **Cluster Mode**: Driver is inside AM in the cluster
18. **Types of transformation**: Narrow and Wide
19. **Narrow**: Data shuffling doesn't happen (map, flatMap, filter)
20. **Wide**: Data shuffling happens (reduceByKey, groupByKey)
21. **reduceByKey()** is a transformation and **reduce()** is an action
22. **reduceByKey()**: Data is processed at each partition, **groupByKey()** : Data is grouped at each partition and complete processing is done at reducer.
23. **Repartition**: used to increase/decrease partitions. Use it for INCREASE
Coalesce: used to decrease partitions and optimized as data shuffling is less

SPARK DATAFRAMES:

1. **Cache()** : It is used to cache the data only in memory.
`Rdd.cache()`
2. **Persist()** : it is used to cache the data in different storage levels (memory, disc, memory & disc, off heap).
`Rdd.persist(StorageLevel._____)`
3. **Serialization**: Process of converting data in object form into bytes, occupies less space

4. **De-Serialization**: Process of converting data in bytes back to objects, occupies more space.
5. **DAG** : Created when an action is called, represents tasks, stages of a job
6. **Map** : performs one-to-one mapping on each line of input
7. **mapPartitions**: performs map function only once on each partition
8. **Driver**: converts high level programming constructs to low level to be fed to executors (dataframe to rdd)
9. **Executors**: Present in memory to process the rdd
10. **Spark context**: creates entry point into spark cluster for spark appl
11. **Spark session**: creates unified entry point into spark cluster
12. **Data frame**: it is a dataset[row] where type error caught only at run time
13. **Data set**: it is a dataset[object] where type error caught at compile time
14. Modes of dealing with corrupted record: permissive, malformed, fail fast
15. Schema types: implicit, infer, explicit (case class, StructType, DDL string)

SPARK OPTIMIZATIONS

1. Spark optimization:
 - a. Cluster Configuration : To configure resources to the cluster so that spark jobs can process well.

- b. Code configuration: To apply optimization techniques at code level so that processing will be fast.
- 2. **Thin executor**: More no. of executors with less no. of resources. Multithreading not possible, too many broadcast variables required. Ex. 1 executor with each 2 cpu cores, 1 GB ram.
- 3. Fat executor: Less no. of executors with more amount of resources. System performance drops down, garbage collection takes time. Ex 1 executor 16 cpu cores, 32 GB ram.
- 4. Garbage collection: To remove unused objects from memory.
- 5. Off heap memory: Memory stored outside of executors/ jvm. It takes less time to clean objects than garbage collector, used for java overheads (extra memory which directly doesn't add to performance but required by system to carry out its operation)
- 6. Static allocation: Resources are fixed at first and will remain the same till the job ends.
- 7. Dynamic Allocation: Resources are allocated dynamically based on the job requirement and released during job stages if they are no longer required.
- 8. Edge node: It is also called as gateway node which is can be accessed by client to enter into hadoop cluster and access name node.
- 9. How to increase parallelism :
 - a. Salting : To increase no. of distinct keys so that work can be distributed across many tasks which in turn increase parallelism.
 - b. Increase no. of shuffle partitions
 - c. Increase the resources of the cluster (more cpu cores)
- 10. Execution memory : To perform computations like shuffle, sort, join
- 11. Storage memory : To store the cache
- 12. User memory : To store user's data structures, meta data etc.
- 13. Reserved memory : To run the executors

14. **Kyro Serializer**: Used to store the data in disk in serialized manner which occupies less space.
15. **Broadcast join**: Used to send the copies of data to all executors. Used when we have only 1 big table.
16. **Optimization on using coalesce()** rather than repartition while reducing no. of partitions
17. **Join optimizations**:
 - a. To avoid or minimize shuffling of data
 - b. To increase parallelism
 1. How to avoid/minimize shuffling?
 - a. Filter and aggregate data before shuffling
 - b. Use optimization methods which require less shuffling (coalesce())
18. **How to increase parallelism ?**
 - a. **Min (total cpu cores, total shuffle partitions, total distinct keys)**
 - b. Use salting to increase no. of distinct keys
 - c. Increase default no. of shuffle partitions
 - d. Increase resources to inc total cpu cores
19. **Skew partitions**: Partitions in which data is unevenly distributed. Bucketing, partitioning, salting can be used to handle it.
20. **Sort aggregate**: Data is sorted based on keys and then aggregated. More processing time
21. **Hash aggregate**: Hash table is created and similar keys are added to the same hash value. Less processing time.
22. **Stages of execution plan :**
 - a. Parsed logical plan (unresolved logical plan) : To find out syntax errors

- b. **Analytical logical plan (Resolved logical plan)**: Checks for column and table names from the catalog.
- c. **Optimized logical plan (Catalyst optimization)**: Optimization done based on built in rules.
- d. **Physical plan**: Actual execution plan is selected based on cost effective model.
- e. **Conversion into Rdd**: Converted into rdd and sent to executors for processing.

****Note:**

1 hdfs block = 1 rdd partition = 128mb

1 hdfs block in local=1 rdd partition in local spark cluster= 32mb

1 rdd ~ can have n partitions in it

1 cluster = 1 machine

N cores = N blocks can run in parallel in each cluster/machine

N stages = N - 1 wide transformations

N tasks in each stage= N partitions in each stage for that rdd/data frame